Building Scalable Deep Learning Systems with RPC

Jeff Hajewski
Department of Computer Science
University of Iowa
Iowa City, IA, USA
jeffrey-hajewski@uiowa.edu

Suely Oliveira
Department of Computer Science
University of Iowa
Iowa City, IA, USA
suely-oliveira@uiowa.edu

Abstract—word word

Index Terms—distributed deep learning, neural architecture search, artificial intelligence

I. Introduction

Building robust and scalable distributed applications is challenging — getting communications patterns correct, handling node failures, and allowing for elastic compute resources all contribute to a high level of complexity. These challenges are exacerbated for long-runnning applications such as training large deep learning models or neural architecture search. Among the many popular frameworks for distributed programming, MPI [11] combined with additional accelerator paradigms such as OpenMP [9], OpenACC [35], and CUDA [21] is the common choice for performance-critical numerical workloads. In the deep learning setting, MPI is less popular with many favoring mutli-threading in combination with multiple GPUs, and more recently experimenting with Kubernetes [29], [30]. Although part of this is simplicity of communication patterns, the fault tolerance required for long-running model training (which can last on the order of months in the industrial setting) makes MPI a poor choice for the underlying communication infrastructure.

In this work we propose an RPC-based system for distributed deep learning. We experiment with the proposed architecture in the domain of neural architecture search (NAS), a computationally intensive problem that trains thousands of deep neural networks in search of an optimal network architecture. We use this problem domain to illustrate four primary advantages to building a computationally intensive distributed system using RPC:

- RPC's higher level of abstraction over MPI simplifies the process of building more complex systems and communication patterns.
- The user can avoid lower-level message serialization and deserialization through RPC coupled with framework such as Protocol Buffers or Thrift.
- RPC systems don't require underlying software or resource managers — a user can create an AWS instance and immediately run their RPC-based code.
- RPC-based systems can offer elastic compute abilities, adding or removing nodes as needed without needing to stop or restart the system. This can be particularly useful when combined with an orchestration technology such as Kubernetes, Amazon Container Service, Azure Container Service, or Google Container Engine.

Or system consists of four separate pieces: a *model* that directs the search for a network architecture, a number of *workers* that perform the computational work of training the models, a number of *brokers* that form the backbone of the data pipeline from model to workers, and a *nameserver* that simplifies the process of adding new brokers, workers, or models to the system. It also offers elastic compute resources, allowing an arbitrary number of workers to join during high computational loads as well as allowing workers to leave the system, decreasing the overall available compute, without needing to restart or manual intervention. The system is fault-tolerant to the loss of workers or brokers, and is highly scalable due to the ability of the brokers to share work and compute resources. Perhaps most importantly from a usability perspective, our system is language agnostic. In our experiments, we use Python for our model and workers, which we use to build and train or deep neural networks via PyTorch [22], and use Go to build the data pipeline of brokers. We use gRPC [34] to handle the generation of RPC stubs, but could have just as easily used Apache Thrift [2], which generates stubs in a larger range of languages such as Ocaml, Haskell, and Rust.

Although we are proposing RPC as an alternative to MPI, it is important to note that we are not claiming RPC is superior to MPI in every problem domain. In the domain of distributed deep learning, RPC offers many useful features such as simple fault tolerance and allowing compute resources to join and leave the system at will. In other domains, such as distributed linear algebra computations, MPI is probably a better choice due to its built in ability to efficiently broadcast messages to nodes within a system in a manner that takes advantage of the network architecture.

II. Neural Architecture Search

The goal of neural architecture search (NAS) is to find an optimal neural network architecture for a given problem.

Fig. 1. Reduction layer.

NAS is computationally intensive due to the requirement of having to train a candidate network in order to evaluate its effectiveness. Although recent novel approaches have dramatically reduced this cost [6], [23], these techniques fix certain elements of the design process, somewhat limiting the available architectures. Despite the computationally intense nature of the NAS problem, the task itself is trivially parallelizable across the network evaluations — two separate networks can be trained simultaneously before being evaluated against eachother.

The two common approaches to NAS are reinforcement learning based approaches such as [16], [23], [37], and evolutionary approaches such as [17], [19], [25]. In our experiments we focus on the evolutionary approach due to its simplicity both in understanding and implementation.

A. Evolutionary Algorithms for NAS

We use a relatively simple approach to evolving neural network architectures in that we only construct linear networks, rather than allow an arbitrary number of incoming connections for any given layer. The motivation behind this is two-fold: it reduces the size of search space for the neural network architecture and simplifies the implementation. Because this the focus of this work is the system architecture, rather than neural architecture search, we feel the trade-off is reasonable.

The problem domain we focus on is computer vision, so we restrict ourselves to convolutional layers for the hidden layers. The networks are constructed by duplicating modules one after the other a pre-determined (by the user) number of times, where layers composing the module unit are determined via evolutionary search. This approach allows us to control the decrease in spatial dimensions, which is useful in the implementation of the network (all layers within a module have the same spatial dimension) and also allows us to force the heuristic of doubling the number of filters via a 1x1 convultional layer prior to reducing the spatial dimension by a factor of two via a 3x3 max pooling layer. The intuition behind this technique is to try and preserve as much information contained within a tensor as possible.

Module evolution proceeds by appending layers to the module or performing crossover, where two network are split at random positions and recombined to form two new networks. The maximum number of layers in a module is fixed (by the user), and evolution can add any of the following 2D convolutions:

- 1x1 Convolution
- 3x3 Convolution
- 3x1-1x3 Convolution
- 5x5 Convolution
- 5x1-1x5 Convolution

A ReLU activation is used between layers

Algorithm 1 High-level outline of evolutionary algorithm.

1: **procedure** PRODUCEOFFSPRINT(N_1, N_2)
2:     **if** $|N_1| \neq |N_2|$ **then**
3:         $o_1 \leftarrow N_1.$mutate()
4:         $o_2 \leftarrow N_2.$mutate()
5:         **return** ReturnFittest(o_1, o_2)
6:     **end if**
7:     $o \leftarrow N_1.$crossover(N_2)
8:     **return** $o.$mutate()
9: **end procedure**
10: **procedure** MUTATE
11:     *// Possibly append layer*
12:     **if** sampleUniform(0, 1) $<$ append_p **then**
13:         self.layers.append(randomLayer())
14:     **end if**
15: **end procedure**

III. RPC-BASED COMMUNICATION

Remote Procedure Call (RPC) offers a method of invoking a function on a remote computer with a given set of arguments. Most RPC frameworks involve a DSL used to define the RPC service, that is, the API available to the caller, and some type of data serialization format. For example, gRPC uses Protocol Buffers (ProtoBufs) [32] as the serialization format for data sent across the network. RPC offers a number of advantages for network communication. It is robust to node failures or network partitions (the RPC invokation simply fails). The data sent across the network is compactly represented, giving way to high bandwidth and low latency communication. The point-to-point communication allows for diverse communication patterns and paradigms. RPC forms the network communication infrastructure at Google [31], Facebook [2], as well as Hadoop [18], [28].

RPC frameworks typically use an IDL to represent the RPC service. Figure 3 gives an example of a gRPC service definition for a heartbeat service. The ProtoBuf compiler uses this definition to generate server stubs and service clients in a number of languages (C++, Java, Python, Go, etc.). The receiver of the RPC call must complete the server stubs by implementing the defined interface. For example, in the heartbeat service defined in Figure 3, the receiver of the heartbeat message would implement a SendHeartbeat function whose body would handle the logic of *receiving* a heartbeat from another process, such as updating a timestamp for the given process ID. The caller of SendHeartbeat uses the generated Heartbeat service client and is only responsible for constructing the HeartbeatMsg.

While the robustness to node failures and finer-grained point-to-point communication capabilities of RPC are core to building resilient distributed systems, the more powerful
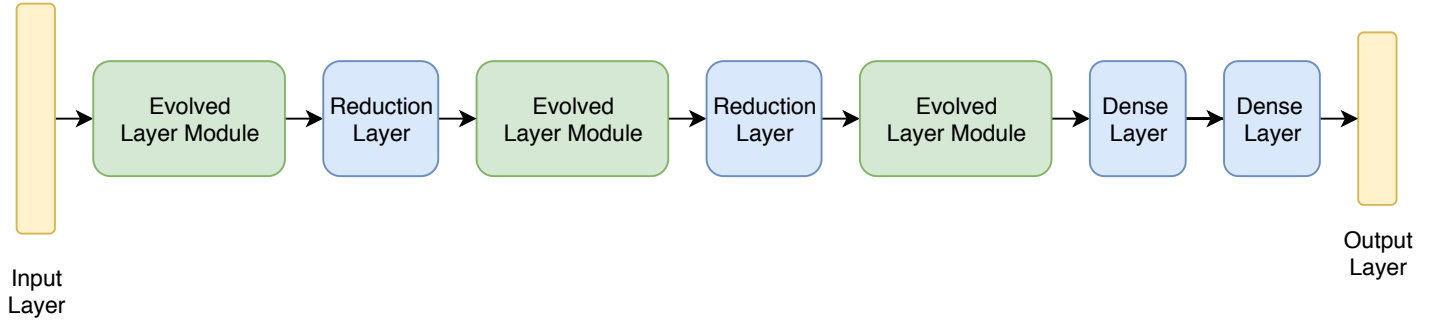
Fig. 2. Example of a constructed network.

```
1  service Heartbeat {
2    rpc SendHeartbeat(Heartbeat) returns (HBResp) {}
3  }
4  message Heartbeat {
5    string id = 1;
6  }
7  message HBResp {}
```

Fig. 3. Example gRPC service definition of a heartbeat service.

```
1  message Task {
2    string id = 1;
3    enum type = 2; // or string type
4    bytes task_obj = 3;
5  }
```

Fig. 4. Example of a data agnostic message type.

```
1   service Task {
2     // Called by a worker to request a task.
3     rpc RequestTask(TaskRequest)
4       returns (TaskResponse) {}
5   }
6   service Result {
7     // Called by a worker to send a result
8     // back to the broker.
9     rpc SendResult(ResultMsg)
10      returns (ResultResponse) {}
11  }
```

Fig. 5. Example service definitions for tasks and results.

RPC API of the system is the service definition itself acts as documentation on the flow of information within the system. A user can look at these definitions and see how information is meant to flow.

feature we capitalize on is the ability to build a system that is agnostic to the type of data flowing through its pipes. Figure 4 gives an example of constructing a Protocol Buffer[CITE] message type that can be used to transport arbitrary data types through the system. Protocol Buffers (and similarly, Thrift messages) support arbitrary length byte arrays [1] This means the user can send a serialized object stored in the `Task`'s `task_obj` field without having to modify the system. A user can switch between running models and tasks using Java to running models and tasks using Python without modifying the system. In fact, the system can transport and run these tasks (in both Java and Python) simultaneously. By encapsulating the tasks (and results) as serialized objects stored as byte arrays the entire system can be data type agnostic.

There is a slight catch. If information within the serialized object is needed to properly schedule or transport the task/result to its destination, this information will need to be added to the message definition. This requires recompiling the message types and regenerating the gRPC (or Thrift) stubs. This is a minor inconveninece, as incorporating this new information into the system requires modifying the system infrastructure.

Another advantage of using gRPC or Thrift to define the

A. RPC vs MPI

From a performance perspective, some prior work [24] has found RPC to provide lower latency and higher bandwidth for some tasks. That does not tell the entire story. MPI comes with built-in complex communication primitives that are capable of taking advantage the physical network architecture of the cluster. Additionally, MPI is capable of bulk broadcasts to groups of nodes using custom communicator definitions. These primitives save the programmer a lot of overhead when building HPC applications. Settings such as large, distributed matrix multiplications or iterative optimization algorithms are particularly well suited for MPI's communication primitives.

However, applications that benefit from specialized communication patterns or that are asynchronous in nature are typically better suited for RPC-based communication. Consider, for example, an application that sends some amount of work to a number of worker nodes and waits for the work to be completed. In the synchronous setting, the entire system will only be as fast as its slowest node and will cause idle resources as the faster nodes wait for the slowest node to finish. In the asynchronous setting, the faster nodes can continue to process work. This scenario was a motivation for [10].

Another short-coming of RPC is message size limitations. This can be particularly troublesome in distributed machine

---

[1]In practice these arrays are limited in size. The failure rate of the RPC system typically increases with the size of the messages.

learning settings where highly parameterized models such as deep neural networks are sent across the network. The solution is to send a representation of the model, rahter than the model itself, across the network. The advantage of this approach is reduced network bandwidth utilization. The disadvantage of course is the loss of trained models at the worker nodes. In practice, this may not be a major issue as final model architectures are typically trained in a specialized manner (e.g., for a large number of epochs). Additionally, models can be saved to a distributed filesystem such as HDFS.

Of course, MPI does offer asynchronous communication, but it is not as flexible as that of RPC-based systems. In an asynchronous setting there are usually classes of nodes whose behavior are class-dependent. In the system we propose, there are four classes of processes: the model, the broker, the nameserver, and the worker. It is conceptually simpler to think of these as four different processes/programs and build each of them separately, rather than a single program that determines its behavoir based on its assigned communicator rank (as done in MPI).

An interesting advantage of RPC is the ability to take advantage of a container orchestration system such as Kubernetes [26] to make sure a specified number of worker nodes remain available. If a worker node goes down, or the desired number of worker nodes increases, Kubernetes can handle the logistics and the system takes advantage of the newly available or restored compute resources without issue. This is only possible because RPC allows us to create a stateless communication that doesn't require a persistent connection between communicating nodes. While this may seem excessive for a small collection of nodes, the ability to have Kubernetes manage system state within a large system can be quite valuable in a production system.

One particularly unique aspect of the domain of neural architecture search is that it does not require low-latency communication. Each neural network can take anywhere from minutes to tens of minutes to train, which means the frequency of task and corresponding result messages is fairly low. One prior work [25] even used a shared filesystem to communicate between a model and the worker nodes. The long time interval between sending the task to the worker and receiving the result from the same worker means a single server could handle a potentially large number of connections with workers, certainly on the order of hundreds and possibly even thousancs. Anecdotally, we found a laptop with an Intel 4.2GHz i7 processor was able to handle over 200 connections of a heartbeat request and response service, similar to that depicted in Figure 3.

IV. SYSTEM ARCHITECTURE

As previously mentioned, our system consists of four different components. Figure 6 shows an overview of the system architecture. A model is a problem specific implementation that controls what is sent to the system for evaluation and handles the result it receives. The brokers form the data pipeline of the system, moving work to available workers.


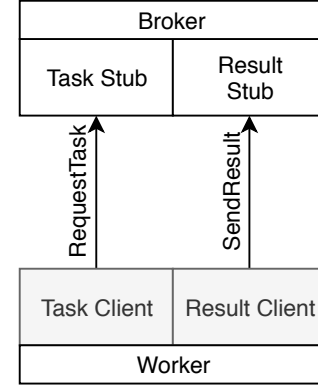Fig. 6. Diagram of system architecture.

Fig. 7. Communication pattern between a worker and broker.

Workers form the other customizable part of the system because they need to know how to perform their assigned work. Lastly, the nameserver maps known brokers to their network address – this is useful for connecting to brokers, such as a model connecting to a broker, a broker connecting to a broker (for broker-broker peering), or a worker connecting to a broker. The following sections will detail each components functionality individual and within the system as a whole.

A. Model

The model is the problem-specific, user-defined logic that determines what work should be performed next and how the results of previously assigned work should be processed. The only requirement of the model is that is uses a broker client stub (generated by gRPC) to push work to the system and implements the model service interface to allow the broker to push results back to the model.

The model needs to track outstanding tasks that have been sent to the broker. While the system is fault-tolerant for most brokers and all workers, if the broker the model is sending work to fails, the work the model is waiting to receive will be lost and the model will need to resend to a new broker.

B. Worker

Workers are the other user-defined and implemented portion of the system. While a single worker implementation can work for multiple model implementations, there is no general worker implementation that will work across all languages and models.

```
1  class BaseTask:
2    def run(self):
3      raise NotImplementedError()
4
5  class Worker:
6    def process_task(self, task):
7      result = task.run()
8      self.broker_client.send_result(result)
```

Fig. 8. Example Task API in Python.

Using an API similar to that shown in Figure 8, one can use the same worker implementation for any task that inherits from the `BaseTask` class.

Figure 7 shows the communication pattern between the broker and a worker. There are two interesting aspects of this diagram: all communication starts at the worker and there is no heartbeat exchange between the worker and the broker. Both of these details are what allow the system to treat workers as ephemeral compute resources. This means an arbitrary number of workers can join and leave the system without the system knowing or caring. The downside to this approach is it requires a little extra bookkeeping at the broker. The broker needs to track what tasks are outstanding and decide if they should be moved back in to the work queue.

C. Broker

Brokers form the data pipeline of our system. Work is sent from a model to a broker, which in turn sends the work to a free worker or to another broker via peering and returns the result to the original model. At its core, a broker is essentially just a process with a owned task queue, a helper task queue (tasks received from other brokers via peering), a processing queue, and a results queue. Work in the owned task queue is work that was sent from a model directly to the broker – this is the work that will be lost if the broker crashes. Work in the helper task queue is work that has been sent from other brokers that the respective broker has agreed to help with. If the broker crashes, this work will *not* be lost as the other brokers will see the failure and can recover the task from their processing queue. The processing queue stores tasks that have been sent to workers or other brokers. When a result is received from a worker or another broker, the corresponding ID will be removed from the processing queue and the result will be added to the result queue. The broker pulls tasks from the result queue and sends the result to its owner, which may be a model or another broker.

D. Nameserver

The nameserver simplifies bookkeeping when starting new broker instances. Rather than forcing the user to specify which brokers a newly started broker should link up with, the nameserver stores and shares that information with all registered brokers. During start-up, each broker registers with the nameserver and begins sending heartbeat messages. The nameserver tracks which brokers have sent heartbeats recently
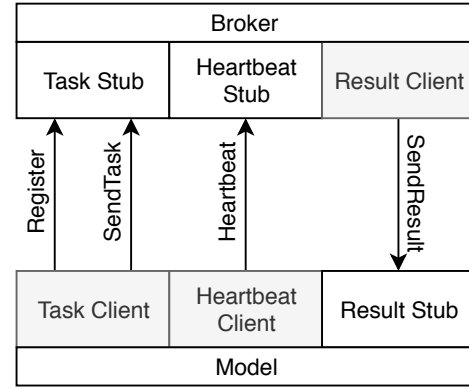

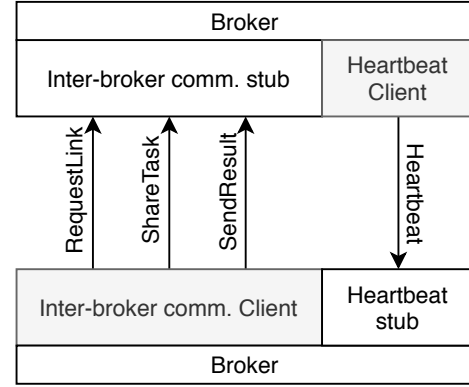Fig. 9. Communication pattern between the broker and model.

Fig. 10. Communication pattern for broker-broker communication.

(via a user-modifiable timeout setting) and drops brokers that have timed-out. If a broker sends a heartbeat *after* the name-server has dropped the connection, the nameserver responds by telling the broker it must re-register with the nameserver.

The nameserver simplifies broker-broker peering by providing a central location where brokers can request addresses of other brokers. It also provides a metadata store for the system, which is important in multi-GPU scenarios. A worker running on a server with multiple GPUs needs information about what GPU to use. Storing this metadata on the nameserver allows the nameserver to manage GPU resources, rather than requiring the user to specify the GPU ID for every worker that is running. This is not a problem for small numbers of workers (<10) but is not scalable for hundreds or thousands of workers.

V. RESULTS

The goal of our experiments is to explore the scalability and robustness of a system built using RPCs for communication to perform a long-running, computationally intensive task. All experiment were run on Amazon's AWS EC2 platform. The first run on a single machine sets the baseline number of models evaluated per unit of time. The successive experiments analyze how this value changes with the addition of worker nodes as well as how the system reacts to losing worker nodes.

Fig. 11. Validation accuracy over time on a single Nvidia V100 for the CIFAR10 dataset.

Figure 11 shows the results of a single Nvidia V100 graphics card. This run establishes a baseline of a single machine performing a naive neural architecture search.

VI. RELATED WORK

The idea of a brokered message queue is not new. RabbitMQ [CITE] is a general purpose message broker that supports the same functionality demonstrated in this paper, but messages are delivered based on a routing key, rather than the first available worker. Similar to the RPC message definitions used in this work, RabbitMQ uses a collection of bytes as the message body. For streaming data, Apache Kafka [15] is a good choice. Kafka is requires a ZooKeeper [12] instance and is generally more complex to setup and run. Both RabbitMQ and Kafka are robust to failures. At the other end of the spectrum, ∅MQ is a low-level messaging library that can be used to build a performant, brokered messaging system similar to the one described in this paper.

A natural question at this point in the work is why we would bother building our own brokered system when there are industrial strength alternatives available. These message passing systems[2] stricly deal with sending and delivering messages in a scalable and robust manner. Our broker does more than simply relay messages—the broker is determining retry logic, hadling lost tasks, [ADD OTHER STUFF].

Tensorflow [1] offers some distributed training facilities. Other work such as Horovod [27] has improved upon Tensorflow's built-in distributed training facilities. Work has also been done on training deep learning models using MPI [8], [33]. However, all of these works suffer from fault tolerance of workers and elastic compute environments. Additionally, they require that MPI is installed on the cluster, which is not an issue if the software is running on a University research cluster or at a national lab, but requires the user to setup the underlying MPI installation if run on a provisioned AWS, GCP, or Azure cluster.

[2]We use this as a generic term, since Kafka is not really a message queue.

MXNet [7] introduces a declarative language used to build a computation graph, which is then optimized and evaluated by the MXNet engine, similar to Tensorflow pre-2.0. [NEEDS ADDITIONAL WORK]

Much of the previous work on neural architecture search uses some form of a distributed architecture consisting of a model, (possible) a coordinator, and workers. The coordinator handles assignment of work to the workers. While many of these works don't detail the specifics of their system, we reviewed some of the available code on Github. A popular paradigm is using Python's `multiprocessing` module to run multiple models on a multi-GPU machine. These GPUs are fed from a thread-safe queue. PyTorch offers a data parallel module that handles this functionality but currently struggles to fully utilize all available GPUs in settings with 8+ GPUs, per the PyTorch documentation.Tensorflow provides similar functionality for distributed training but instead relies on gRPC.

A number of prior works use MPI in conjunction with deep learning libraries taking a data parallel approach [4], [5]. This is familiar to other work that takes advantage of multithreading (such as Python's `multiprocessing` module) but utilizes much larger systems. Awan et. al. [3] build a model parallel system using MPI, relying on MPI's efficient communication primitives to avoid excessive blocking when cross-node dependencies. Unfortunately, all of these works require MPI on the system they are running, which adds another dependency to the software stack. These works also are unable to utilize newly available resources without restarting the running process.

Others [14], [20] use Spark [36] to train deep learning neural networks. Spark *can* handle the addition or loss of compute nodes, is robust to node failures (via HDFS) but struggles with iterative algorithms. Both of these works use Spark to distribute Caffe [13] models to GPU compute nodes and implement asynchronous SGD. The main issue with Spark in this context is that either the user must create a wrapper for the model in Scala or use PySpark. Using PySpark is relatively simple, as long as your model is written in Python. If your model is written in something other than Python such as C++, then you must decide if its better to create a wrapper in Scala (if you know Scala) or a Swig interface and call it from Python. Both of these options seem more complex than implementing a server stub to send models and using a pre-generated client to request models to train as well as report the results.

Implementing algorithms that are iterative in nature in Spark poses another issue. Spark does not handle iterative-style algorithms particularly well, primarily due to the shuffle-stage of MapReduce (since Spark is built on top of Hadoop).

VII. CONCLUSION

It can be tempting to think of system design as an all-or-nothing decision—either build a system with MPI or build a system using RPCs. An interesting avenue for future work is combining the two. Consider a workload whose core unit of work is amenable to an MPI-based system but the individual units of work are independent of eachother with the exception

of possible boundary interactions. A combination of RPC and MPI communication might be ideal—using MPI within a single cluster to complete individual units of work and RPC to communicate between clusters. In this way, individual clusters become the workers of the system and can join and leave at will, while the broker backbone manages sending tasks to the individual clusters and returning their results to the model.

REFERENCES

[1] M. Abadi, A. Agarwal, P. Barham, E. Brevdo, Z. Chen, C. Citro, G. S. Corrado, A. Davis, J. Dean, M. Devin, S. Ghemawat, I. Goodfellow, A. Harp, G. Irving, M. Isard, Y. Jia, R. Jozefowicz, L. Kaiser, M. Kudlur, J. Levenberg, D. Mané, R. Monga, S. Moore, D. Murray, C. Olah, M. Schuster, J. Shlens, B. Steiner, I. Sutskever, K. Talwar, P. Tucker, V. Vanhoucke, V. Vasudevan, F. Viégas, O. Vinyals, P. Warden, M. Wattenberg, M. Wicke, Y. Yu, and X. Zheng. TensorFlow: Large-scale machine learning on heterogeneous systems, 2015. Software available from tensorflow.org.

[2] A. Agarwal, M. Slee, and M. Kwiatkowski. Thrift: Scalable cross-language services implementation. Technical report, Facebook, 4 2007.

[3] A. A. Awan, C.-H. Chu, H. Subramoni, and D. K. Panda. Optimized broadcast for deep learning workloads on dense-gpu infiniband clusters: Mpi or nccl? In *Proceedings of the 25th European MPI Users' Group Meeting*, EuroMPI'18, pages 2:1–2:9, New York, NY, USA, 2018. ACM.

[4] A. A. Awan, K. Hamidouche, J. M. Hashmi, and D. K. Panda. S-caffe: Co-designing mpi runtimes and caffe for scalable deep learning on modern gpu clusters. In *Proceedings of the 22Nd ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming*, PPoPP '17, pages 193–205, New York, NY, USA, 2017. ACM.

[5] A. A. Awan, K. Hamidouche, A. Venkatesh, and D. K. Panda. Efficient large message broadcast using nccl and cuda-aware mpi for deep learning. In *Proceedings of the 23rd European MPI Users' Group Meeting*, EuroMPI 2016, pages 15–22, New York, NY, USA, 2016. ACM.

[6] A. Brock, T. Lim, J. M. Ritchie, and N. Weston. SMASH: one-shot model architecture search through hypernetworks. *CoRR*, abs/1708.05344, 2017.

[7] T. Chen, M. Li, Y. Li, M. Lin, N. Wang, M. Wang, T. Xiao, B. Xu, C. Zhang, and Z. Zhang. Mxnet: A flexible and efficient machine learning library for heterogeneous distributed systems. *CoRR*, abs/1512.01274, 2015.

[8] A. Coates, B. Huval, T. Wang, D. Wu, B. Catanzaro, and N. Andrew. Deep learning with cots hpc systems. In S. Dasgupta and D. McAllester, editors, *Proceedings of the 30th International Conference on Machine Learning*, volume 28 of *Proceedings of Machine Learning Research*, pages 1337–1345, Atlanta, Georgia, USA, 17–19 Jun 2013. PMLR.

[9] L. Dagum and R. Menon. Openmp: An industry-standard api for shared-memory programming. *IEEE Comput. Sci. Eng.*, 5(1):46–55, Jan. 1998.

[10] J. Dean, G. S. Corrado, R. Monga, K. Chen, M. Devin, Q. V. Le, M. Z. Mao, M. Ranzato, A. Senior, P. Tucker, K. Yang, and A. Y. Ng. Large scale distributed deep networks. In *NIPS*, 2012.

[11] M. P. Forum. Mpi: A message-passing interface standard. Technical report, Knoxville, TN, USA, 1994.

[12] P. Hunt, M. Konar, F. P. Junqueira, and B. Reed. Zookeeper: Wait-free coordination for internet-scale systems. In *Proceedings of the 2010 USENIX Conference on USENIX Annual Technical Conference*, USENIXATC'10, pages 11–11, Berkeley, CA, USA, 2010. USENIX Association.

[13] Y. Jia, E. Shelhamer, J. Donahue, S. Karayev, J. Long, R. Girshick, S. Guadarrama, and T. Darrell. Caffe: Convolutional architecture for fast feature embedding. In *Proceedings of the 22Nd ACM International Conference on Multimedia*, MM '14, pages 675–678, New York, NY, USA, 2014. ACM.

[14] H. Kim, J. Park, J. Jang, and S. Yoon. Deepspark: Spark-based deep learning supporting asynchronous updates and caffe compatibility. *CoRR*, abs/1602.08191, 2016.

[15] J. Kreps, N. Narkhede, and J. Rao. Kafka: a distributed messaging system for log processing. In *NetDB*, 2011.

[16] G. Kyriakides and K. G. Margaritis. Neural architecture search with synchronous advantage actor-critic methods and partial training. In *Proceedings of the 10th Hellenic Conference on Artificial Intelligence*, SETN '18, pages 34:1–34:7, New York, NY, USA, 2018. ACM.

[17] H. Liu, K. Simonyan, O. Vinyals, C. Fernando, and K. Kavukcuoglu. Hierarchical representations for efficient architecture search. *CoRR*, abs/1711.00436, 2017.

[18] X. Lu, N. S. Islam, M. Wasi-Ur-Rahman, J. Jose, H. Subramoni, H. Wang, and D. K. Panda. High-performance design of hadoop rpc with rdma over infiniband. In *Proceedings of the 2013 42Nd International Conference on Parallel Processing*, ICPP '13, pages 641–650, Washington, DC, USA, 2013. IEEE Computer Society.

[19] R. Miikkulainen, J. Z. Liang, E. Meyerson, A. Rawal, D. Fink, O. Francon, B. Raju, H. Shahrzad, A. Navruzyan, N. Duffy, and B. Hodjat. Evolving deep neural networks. *CoRR*, abs/1703.00548, 2017.

[20] P. Moritz, R. Nishihara, I. Stoica, and M. I. Jordan. Sparknet: Training deep networks in spark. In Y. Bengio and Y. LeCun, editors, *ICLR (Poster)*, 2016.

[21] J. Nickolls, I. Buck, M. Garland, and K. Skadron. Scalable parallel programming with cuda. *Queue*, 6(2):40–53, Mar. 2008.

[22] A. Paszke, S. Gross, S. Chintala, G. Chanan, E. Yang, Z. DeVito, Z. Lin, A. Desmaison, L. Antiga, and A. Lerer. Automatic differentiation in PyTorch. In *NIPS Autodiff Workshop*, 2017.

[23] H. Pham, M. Guan, B. Zoph, Q. Le, and J. Dean. Efficient neural architecture search via parameters sharing. In J. Dy and A. Krause, editors, *Proceedings of the 35th International Conference on Machine Learning*, volume 80 of *Proceedings of Machine Learning Research*, pages 4095–4104, Stockholmsmässan, Stockholm Sweden, 10–15 Jul 2018. PMLR.

[24] K. Qureshi and H. Rashid. A performance evaluation of rpc, java rmi, mpi and pvm. *Malaysian Journal of Computer Science*, 18:38–44, 01 2006.

[25] E. Real, S. Moore, A. Selle, S. Saxena, Y. L. Suematsu, J. Tan, Q. V. Le, and A. Kurakin. Large-scale evolution of image classifiers. In D. Precup and Y. W. Teh, editors, *Proceedings of the 34th International Conference on Machine Learning, ICML 2017, Sydney, NSW, Australia, 6-11 August 2017*, volume 70 of *Proceedings of Machine Learning Research*, pages 2902–2911. PMLR, 2017.

[26] D. K. Rensin. *Kubernetes - Scheduling the Future at Cloud Scale*. 1005 Gravenstein Highway North Sebastopol, CA 95472, 2015.

[27] A. Sergeev and M. D. Balso. Horovod: fast and easy distributed deep learning in tensorflow. *CoRR*, abs/1802.05799, 2018.

[28] K. Shvachko, H. Kuang, S. Radia, and R. Chansler. The hadoop distributed file system. In *Proceedings of the 2010 IEEE 26th Symposium on Mass Storage Systems and Technologies (MSST)*, MSST '10, pages 1–10, Washington, DC, USA, 2010. IEEE Computer Society.

[29] S. Song, L. Deng, J. Gong, and H. Luo. Gaia scheduler: A kubernetes-based scheduler framework. In *2018 IEEE Intl Conf on Parallel Distributed Processing with Applications, Ubiquitous Computing Communications, Big Data Cloud Computing, Social Computing Networking, Sustainable Computing Communications (ISPA/IUCC/BDCloud/SocialCom/SustainCom)*, pages 252–259, Dec 2018.

[30] P. Tsai, H. Hong, A. Cheng, and C. Hsu. Distributed analytics in fog computing platforms using tensorflow and kubernetes. In *2017 19th Asia-Pacific Network Operations and Management Symposium (APNOMS)*, pages 145–150, Sep. 2017.

[31] J. Van Winkel and B. Beyer. The production environment at google, from the viewpoint of an sre. *Site Reliability Engineering: How Google Runs Production Systems*, 2017.

[32] K. Varda. Protocol buffers: Google's data interchange format. Technical report, Google, 6 2008.

[33] A. Vishnu, C. Siegel, and J. Daily. Distributed tensorflow with MPI. *CoRR*, abs/1603.02339, 2016.

[34] X. Wang, H. Zhao, and J. Zhu. Grpc: A communication cooperation mechanism in distributed systems. *SIGOPS Oper. Syst. Rev.*, 27(3):75–86, July 1993.

[35] S. Wienke, P. Springer, C. Terboven, and D. an Mey. Openacc: First experiences with real-world applications. In *Proceedings of the 18th International Conference on Parallel Processing*, Euro-Par'12, pages 859–870, Berlin, Heidelberg, 2012. Springer-Verlag.

[36] M. Zaharia, R. S. Xin, P. Wendell, T. Das, M. Armbrust, A. Dave, X. Meng, J. Rosen, S. Venkataraman, M. J. Franklin, A. Ghodsi, J. Gonzalez, S. Shenker, and I. Stoica. Apache spark: A unified engine for big data processing. *Commun. ACM*, 59(11):56–65, Oct. 2016.

[37] B. Zoph and Q. V. Le. Neural architecture search with reinforcement learning. 2017.